

In [3]:

```
import pandas as pd
import numpy as np
import seaborn as sns
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

In [4]:

```
#loading dataset
from sklearn.datasets import load_boston
boston=load_boston()
```

In [5]:

```
boston.keys()
```

Out[5]:

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename'])
```

In [6]:

```
print(boston.DESCR)
```

```
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B $1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/> (<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>)

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression

problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [7]:

```
print(boston.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
 'B' 'LSTAT']
```

In [8]:

```
bs=pd.DataFrame(boston.data)
```

In [9]:

```
bs.head()
```

Out[9]:

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

In [10]:

```
bs.columns=boston.feature_names
```

In [11]:

```
bs.head()
```

Out[11]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	

In [12]:

```
bs['PRICE']=boston.target
```

In [13]:

```
bs['PRICE']
```

Out[13]:

```
0      24.0
1      21.6
2      34.7
3      33.4
4      36.2
...
501     22.4
502     20.6
503     23.9
504     22.0
505     11.9
Name: PRICE, Length: 506, dtype: float64
```

In [14]:

```
bs.head()
```

Out[14]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	

In [15]:

```
bs.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   CRIM        506 non-null    float64
 1   ZN          506 non-null    float64
 2   INDUS       506 non-null    float64
 3   CHAS        506 non-null    float64
 4   NOX         506 non-null    float64
 5   RM          506 non-null    float64
 6   AGE         506 non-null    float64
 7   DIS         506 non-null    float64
 8   RAD         506 non-null    float64
 9   TAX         506 non-null    float64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       506 non-null    float64
13   PRICE       506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

In [16]:

```
bs.isnull().sum()
```

Out[16]:

```
CRIM      0
ZN         0
INDUS      0
CHAS       0
NOX        0
RM         0
AGE        0
DIS        0
RAD        0
TAX        0
PTRATIO    0
B          0
LSTAT      0
PRICE      0
dtype: int64
```

In [17]:

```
bs.describe()
```

Out[17]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	50
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	1

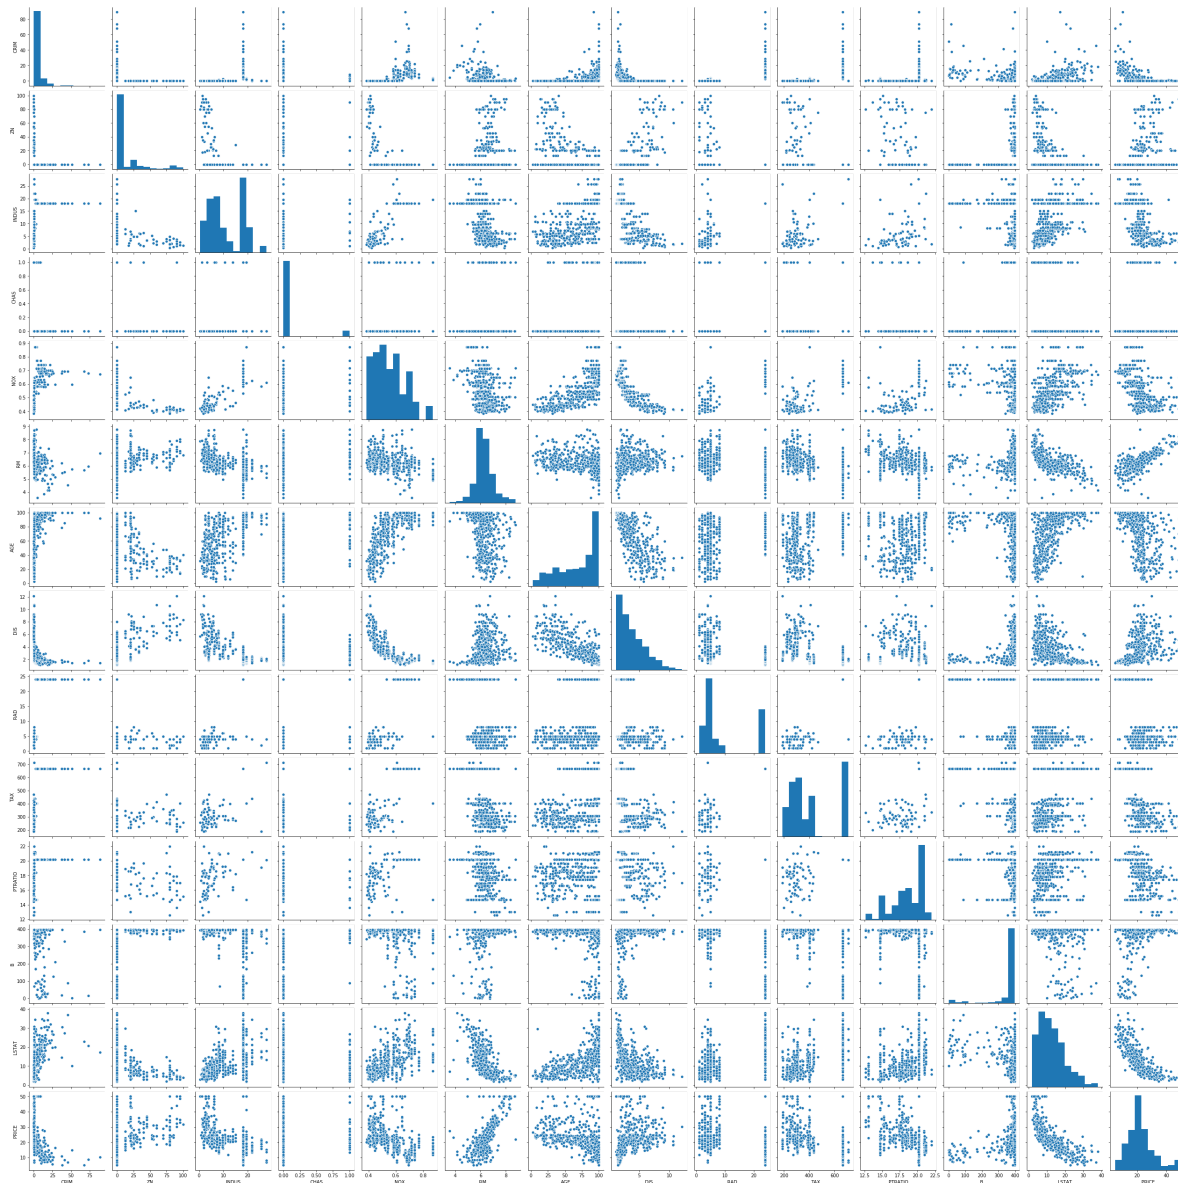


In [18]:

```
sns.pairplot(bs)          #all features are mapped with other features
```

Out[18]:

<seaborn.axisgrid.PairGrid at 0x20e5892dec8>




In [31]:

```
sns.distplot(bs['PRICE'])
```

#visualize the distribution of price

Out[31]:

<matplotlib.axes._subplots.AxesSubplot at 0x20e63cf58c8>




In [22]:

```
sns.lmplot(x='RM',y='PRICE',data=bs)
```

Out[22]:

<seaborn.axisgrid.FacetGrid at 0x20e63c58e08>



Training a linear regression model

In [24]:

```
X=bs.drop('PRICE',axis=1)  
Y=bs['PRICE']
```

Train Test Split

In [25]:

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.3,random_state=4)
```

Creating and training model

In [26]:

```
lr=LinearRegression()
```

In [27]:

```
lr.fit(X_train,Y_train)
```

Out[27]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=False)
```

In [28]:

```
Y_pred=lr.predict(X_test)
```

In [29]:

Y_pred

Out[29]:

```
array([[11.07380893, 26.47910329, 17.34489869, 19.1948608 , 36.3617073
5,
      24.77095832, 31.00051311, 19.94226881, 19.22375105, 24.4299843
5,
      28.31512637, 28.40796034, 19.27427968, 33.82295207, 21.2859648
7,
      15.11171444, 20.97688767, 11.28556596, 11.8611348 , 13.8844438
7,
      5.37422679, 17.55278177, 20.58171204, 22.59849951, 16.0754426
5,
      20.45924503, 19.1068775 , 14.37832191, 21.23235601, 17.5218656
4,
      14.40725559, 23.68483414, 33.7410661 , 22.02733357, 17.6213914
7,
      19.97241153, 30.24069397, 34.69718954, 23.85821534, 24.3071509
3,
      36.13378112, 31.97532293, 19.626175 , 31.61097971, 34.5812780
9,
      25.62718797, 39.95041812, 17.60880538, 19.90319708, 23.4041750
1,
      33.70182396, 25.62491083, 18.25559302, 27.27317174, 13.4637787
1,
      23.43470656, 24.43721849, 33.52056736, 16.99896935, 37.9446440
4,
      15.94567818, 19.32528916, 31.84088262, 15.25081303, 38.4034478
9,
      27.45372884, 34.36154312, 9.37353936, 19.42580066, 21.9921845
9,
      22.79983394, 22.50810313, 22.30918714, 27.84395887, 16.4081834
5,
      22.55507669, 16.5147332 , 25.11106836, 13.76991927, 19.7865639
9,
      22.10247463, 20.26663237, 28.15165586, 19.52050766, 30.3325436
4,
      22.79109999, 29.2663436 , 19.43113706, 24.7968264 , 37.4627564
8,
      31.05503576, 41.3372879 , 18.46365381, 36.67964528, 19.4084240
5,
      23.61810063, 27.93475362, 24.41825213, 9.4599059 , 20.6808867
7,
      8.99426788, 28.4492398 , 31.88237066, 14.04302958, 24.8347909
,
      19.94124425, 36.90271393, 31.06556982, 33.91883403, 28.6459153
6,
      31.1007263 , 22.82363163, 11.58125942, 29.46902405, 37.0606610
6,
      23.01945872, 41.79865192, 18.44334162, 3.433324 , 18.5748566
3,
      22.21257489, 16.71192648, 28.00473344, 28.42374739, 19.6417452
,
      18.76090758, 35.37631447, 13.12349548, 14.73923539, 18.1620233
3,
      38.26604753, 15.97821613, 41.91544265, 30.44631625, 28.6584808
9,
      24.19590457, 12.06559683, 26.01408744, 23.25012698, 18.9250685
```

```
7,
    17.05016777, 17.50245392, 20.89247338, 24.62630514, 1.8216755
8,
    23.03969555, 19.35693345, 17.89193065, 38.43943954, 19.7075262
,
    31.67181183, 19.0130913 ])
```

In []: